JAN FRANCISZEK PIOTROWSKI

janfpiotrowski@gmail.com ◇ +48 884 307 629
jfpio.github.io ◇ github.com/jfpio ◇ linkedin.com/in/jan-piotrowski-23553719b

ABOUT ME

I aim to maximise AI's benefits while reducing systemic risks. My research currently focuses on multi-agent safety and emergent misalignment. I am pursuing a PhD at the Centre for Credible AI, Warsaw University of Technology.

EDUCATION

Centre for Credible AI, Warsaw University of Technology, Poland *2026–present*
PhD Student — research focus: multi-agent safety

University of Warsaw, Poland *2022–2024*
M.Sc. in Machine Learning (GPA 4.58 / 5)
Thesis: Linking news to tweet cascades with a contrastive learning approach

Warsaw University of Technology, Poland *2018–2022*
B.Sc. in Computer Science

PUBLICATIONS

Plisiecki, H., Leniarska, M., **Piotrowski, J.**, Zajenkowski, M. (2026). "Interpretable Semantic Gradients in SSD: A PCA Sweep Approach and a Case Study on AI Discourse." ACL 2026 Findings

Podolak, J., Łukasik, S., Balawender, P., Ossowski, J., **Piotrowski, J.**, Bąkowicz, K., Sankowski, P. (2024). "LLM-generated Responses to Mitigate the Impact of Hate Speech." EMNLP Findings 2024

EXPERIENCE

Centre for Credible AI November 2025–Present
Machine Learning Researcher *Warsaw, PL*

- Investigating whether multi-agent LLM systems deliver genuine capability gains or primarily act as additional test-time compute, using mathematical reasoning tasks as the main evaluation setting.
- Teaching a student research group and mentoring its mini-project on emergent misalignment, which builds on *Model Organisms for Emergent Misalignment* (Turner et al., 2025).
- Member of a red-teaming group on AI control and LLM-based agents; exploring inoculation prompting vs. steering-vector methods and rule-breaking tendencies in LLM agents.

University of Warsaw January 2025 – June 2025
Machine Learning Researcher *Warsaw, PL*

- **RL for Scalable Oversight via LLM Debate.** Extended "Training Language Models to Win Debates with Self-Play Improves Judge Accuracy" to study RL training in multi-agent debate as a scalable-oversight method.

MIM Solutions January 2022 – December 2024
Machine Learning Engineer *Warsaw, PL*

- **Contrastive Learning for Linking Tweets to Events (ERC-funded).** Led a research initiative (supervised by Prof. Sankowski) developing BERT-based dual encoders to link Twitter posts with news articles; owned experimental design, modelling, and analysis. Preprint: arXiv:2312.07599.
- **Automated Counter-Speech to Mitigate Hate Speech (EMNLP Findings 2024).** Took over a stalled student project as lead after two rejections: introduced a novel engagement metric, added ethical analysis, and rewrote the manuscript, leading to acceptance (arXiv:2311.16905).
- **Research Infrastructure.** Trained, evaluated, and fine-tuned deep learning models on GPU clusters using Hydra (config), Optuna (HPO), and Neptune (experiment tracking).

LEADERSHIP

Polish Scouting Association (ZHR) 2017 – Present
Scout Leader, Board Member – Mazovian Region

- Board Member overseeing quality assurance of camps for 2,000+ youth annually; directly managed camps of 120 participants with a 25,000 EUR budget, and mentored 30 leaders.